

PROJECT

Object Classification

A part of the Deep Learning Nanodegree Foundation Program

PROJECT REVIEW	CODE REVIEW	NOTES
----------------	-------------	-------

Requires Changes

SHARE YOUR ACCOMPLISHMENT

3 SPECIFICATIONS REQUIRE CHANGES



Overall great job 🍌 on this resubmission. Just a few small changes and you will be good to go. Please go through the above explanations carefully. Good luck with the Nanodegree!

Required Files and Tests

✓	The project submission contains the project notebook, called "dlnd_image_classification.ipynb".
✓	All the unit tests in project have passed.
	Great job on passing the tests.

Preprocessing

✓	The <code>normalize</code> function normalizes image data in the range of 0 to 1, inclusive.
	Nice job!
✓	The <code>one_hot_encode</code> function encodes labels to one-hot encodings.
	Good job implementing the <code>one_hot_encode</code> function. Here, you could also use the one-liner <code>np.eye(10)[x]</code>

Neural Network Layers


✓	The neural net inputs functions have all returned the correct TF Placeholder.
✓	<p>The <code>conv2d_maxpool</code> function applies convolution and max pooling to a layer.</p> <p>The convolutional layer should use a nonlinear activation.</p> <p>This function shouldn't use any of the tensorflow functions in the <code>tf.contrib</code> or <code>tf.layers</code> namespace.</p> <p>Nice job implementing the <code>conv2d_maxpool</code> function.</p> <p>SUGGESTION</p> <ul style="list-style-type: none"><li>Since you are using ReLU activation, you can save some of the ReLU operations by simply reversing the execution order to Max-Pooling -&gt; ReLU instead of ReLU-&gt;Max-Pooling.</li><li>It is better to set the weights near zero. A small weight leads to a very small gradient, which leads to small updates. Whereas, large weights saturates the activations. Hence it is desirable to keep the weights in such a way that activation functions are in a linear zone so that gradients are optimal. In a deep neural network, a small perturbation in the initial layers leads to a large change in the later layers. Thus, during backpropagation, the gradients have to compensate this, before learning the weights to produce required outputs. Read more about weight initialization <a href="#">here</a>. So it is better to set the standard deviation as something between 0.01 to 0.1(default is 1).</li></ul> <p>Note: This is the main reason why the accuracy is below 50%. I have marked this as correct because conceptually it is correct.</p>
✓	<p>The <code>flatten</code> function flattens a tensor without affecting the batch size.</p> <p>Great job implementing <code>flatten</code> function without using classes from the TensorFlow Layers.</p>
✓	<p>The <code>fully_conn</code> function creates a fully connected layer with a nonlinear activation.</p> <p>Great job implementing <code>fully_conn</code> function without using classes from the TensorFlow Layers. Here also it is better to set the standard deviation to something between 0.01 to 0.1.</p>
✓	<p>The <code>output</code> function creates an output layer with a linear activation.</p> <p>Great job using a linear activation.</p>

Neural Network Architecture

🔄	The <code>conv_net</code> function creates a convolutional model and returns the logits. Dropout should be applied to alt least one layer.
	<p>It is a decent architecture to start with.</p> <p>SUGGESTION</p> <ul style="list-style-type: none"><li>Convolutional network with multiple convolutional layers with increasing depth (e.g. 32, 64, 128), a small convolutional filter (3x3) and stride (1x1) usually work best. Refer <a href="#">here</a> for a detailed discussion and tips for building convolutional network architectures.</li><li>The <code>num_outputs</code> in the Fully Connected Layers is a bit low. Try increasing them along with Dropout and you could easily get &gt;50% validation accuracy! To choose the number of hidden layers and nodes in a feedforward neural network, refer <a href="#">here</a>.</li></ul> <p>Note: The last point is very important in order to get &gt;50% validation accuracy. Also, as per the specification Dropout should be applied to alt least one layer.</p>

Neural Network Training

✓	The <code>train_neural_network</code> function optimizes the neural network.
	Good job! You have correctly used <code>session.run</code>
🔄	The <code>print_stats</code> function prints loss and validation accuracy.
	Good job, you have correctly used the global variables <code>valid_features</code> and <code>valid_labels</code> to calculate validation accuracy. But to calculate the loss you need to use the variables <code>feature_batch</code> and <code>label_batch</code> because here we want to print the loss per batch not the validation loss.
✓	The hyperparameters have been set to reasonable numbers.

	The neural network validation and test accuracy are similar. Their accuracies are greater than 50%.
This will pass when you change the <code>conv2d_maxpool1</code> and <code>conv_net</code> as suggested.	

 RESUBMIT PROJECT

 DOWNLOAD PROJECT



Best practices for your project resubmission

Ben shares 5 helpful tips to get you through revising and resubmitting your project.

 [Watch Video](#) (3:01)

RETURN TO PATH